

BAYESIAN ANALYSIS OF WINE RATINGS

2024-12-08

LOAD IN + UNDERSTANDING THE DATA


```
set.seed(36501)
wine_data<-drop_na(read.csv("winemag-data-130k-v2.csv"))%>%
  ## ONLY USE DATA FROM COUNTRIES WITH OVER 400 DISTINCT REVIEWS
  sample_n(1000)%>%
  filter(country=="Argentina"|country=="Australia"|country=="Austria"|country=="Chile"|country=
="France"|country=="Germany"|country=="Italy"|country=="New Zealand"|country=="Portugal"|country
=="South Africa"|country=="Spain"|country=="US")
wine_data <- wine_data %>%
  mutate(log_price = log(price))
glimpse(wine_data)
```

```
## Rows: 972
## Columns: 15
## $ X                <int> 48954, 22922, 83248, 12855, 99833, 98673, 42224,...
## $ country          <chr> "Portugal", "US", "US", "Australia", "US", "US",...
## $ description      <chr> "This is a finely perfumed wine, full of ripe Ca...
## $ designation      <chr> "Pegos Claros Reserva", "", "", "Phoenix", "West...
## $ points           <int> 92, 88, 88, 91, 89, 82, 88, 81, 84, 90, 88, 89, ...
## $ price            <dbl> 17, 36, 14, 20, 19, 14, 23, 8, 28, 54, 13, 38, 4...
## $ province         <chr> "Palmela", "Oregon", "Oregon", "South Australia"...
## $ region_1         <chr> "", "Willamette Valley", "Willamette Valley", "C...
## $ region_2         <chr> "", "Willamette Valley", "Willamette Valley", ""...
## $ taster_name      <chr> "Roger Voss", "Paul Gregutt", "Paul Gregutt", "J...
## $ taster_twitter_handle <chr> "@vossroger", "@paulgwine ", "@paulgwine ", "@Jo...
## $ title            <chr> "Wines & Winemakers 2011 Pegos Claros Reserva Ca...
## $ variety          <chr> "Castelão", "Pinot Noir", "Riesling", "Cabernet ...
## $ winery           <chr> "Wines & Winemakers", "Trisaetum", "Willamette V...
## $ log_price        <dbl> 2.833213, 3.583519, 2.639057, 2.995732, 2.944439...
```

```
tabyl(wine_data$country)
```

```
## wine_data$country  n    percent
##      Argentina  32 0.032921811
##      Australia  15 0.015432099
##      Austria    21 0.021604938
##      Chile      45 0.046296296
##      France    164 0.168724280
##      Germany    21 0.021604938
##      Italy     133 0.136831276
##      New Zealand  9 0.009259259
##      Portugal   44 0.045267490
##      South Africa 9 0.009259259
##      Spain      47 0.048353909
##      US        432 0.444444444
```

```
ggplot(wine_data,aes(x=log_price))+geom_density()
```



```
mean(wine_data$log_price)
```

```
## [1] 3.313282
```

```
sd(wine_data$log_price)
```

```
## [1] 0.6791097
```

```
wine_data%>%
  group_by(country)%>%
  summarise(mean_points = mean(points, na.rm= TRUE),
            mean_log_price = mean(log_price, na.rm=TRUE),
            n= n())
```

```
## # A tibble: 12 × 4
##   country      mean_points mean_log_price     n
##   <chr>          <dbl>         <dbl> <int>
## 1 Argentina      86.1           2.98     32
## 2 Australia      88.7           3.19     15
## 3 Austria        89.6           3.13     21
## 4 Chile          86.5           2.76     45
## 5 France         89.3           3.46    164
## 6 Germany        90.5           3.44     21
## 7 Italy          88.8           3.40    133
## 8 New Zealand    88.4           3.13      9
## 9 Portugal       88.3           2.93     44
## 10 South Africa  88.2           3.31      9
## 11 Spain         86.9           2.96     47
## 12 US           88.4           3.40    432
```

BAYESIAN ESTIMATION MODEL

BAYESIAN REGRESSION MODEL

```
wine_reg_model<- stan_glmer(
  price ~ points + (points | country),
  data = wine_data, family = gaussian,
  #we know that in this data, the global mean price is around 35, but can be between 22 and 48
  prior_intercept = normal(35,6.5),
  #we believe that a typical country's wine's price will on average be higher when it has a high
er rating, yet we're unsure of the exact rate. We'll say its likely between 1$-5$ per point, how
ever it really could be anything
  prior = normal(3, 1, autoscale =TRUE),
  #we have no idea of the variability between countries or how a country's point ratings might f
luctuate from their trend
  prior_aux = exponential(1, autoscale = TRUE),
  prior_covariance = decov(regularization = 1, concentration = 1, shape = 1, scale = 1),
  chains = 4, iter = 5000*2, seed= 36501)
```

```
##
## SAMPLING FOR MODEL 'continuous' NOW (CHAIN 1).
## Chain 1:
## Chain 1: Gradient evaluation took 0.000183 seconds
## Chain 1: 1000 transitions using 10 leapfrog steps per transition would take 1.83 seconds.
## Chain 1: Adjust your expectations accordingly!
## Chain 1:
## Chain 1:
## Chain 1: Iteration:    1 / 10000 [  0%] (Warmup)
## Chain 1: Iteration: 1000 / 10000 [ 10%] (Warmup)
## Chain 1: Iteration: 2000 / 10000 [ 20%] (Warmup)
## Chain 1: Iteration: 3000 / 10000 [ 30%] (Warmup)
## Chain 1: Iteration: 4000 / 10000 [ 40%] (Warmup)
## Chain 1: Iteration: 5000 / 10000 [ 50%] (Warmup)
## Chain 1: Iteration: 5001 / 10000 [ 50%] (Sampling)
## Chain 1: Iteration: 6000 / 10000 [ 60%] (Sampling)
## Chain 1: Iteration: 7000 / 10000 [ 70%] (Sampling)
## Chain 1: Iteration: 8000 / 10000 [ 80%] (Sampling)
## Chain 1: Iteration: 9000 / 10000 [ 90%] (Sampling)
## Chain 1: Iteration: 10000 / 10000 [100%] (Sampling)
## Chain 1:
## Chain 1: Elapsed Time: 411.795 seconds (Warm-up)
## Chain 1:           245.016 seconds (Sampling)
## Chain 1:           656.811 seconds (Total)
## Chain 1:
##
## SAMPLING FOR MODEL 'continuous' NOW (CHAIN 2).
## Chain 2:
## Chain 2: Gradient evaluation took 0.0001 seconds
## Chain 2: 1000 transitions using 10 leapfrog steps per transition would take 1 seconds.
## Chain 2: Adjust your expectations accordingly!
## Chain 2:
## Chain 2:
## Chain 2: Iteration:    1 / 10000 [  0%] (Warmup)
## Chain 2: Iteration: 1000 / 10000 [ 10%] (Warmup)
## Chain 2: Iteration: 2000 / 10000 [ 20%] (Warmup)
## Chain 2: Iteration: 3000 / 10000 [ 30%] (Warmup)
## Chain 2: Iteration: 4000 / 10000 [ 40%] (Warmup)
## Chain 2: Iteration: 5000 / 10000 [ 50%] (Warmup)
## Chain 2: Iteration: 5001 / 10000 [ 50%] (Sampling)
## Chain 2: Iteration: 6000 / 10000 [ 60%] (Sampling)
## Chain 2: Iteration: 7000 / 10000 [ 70%] (Sampling)
## Chain 2: Iteration: 8000 / 10000 [ 80%] (Sampling)
## Chain 2: Iteration: 9000 / 10000 [ 90%] (Sampling)
## Chain 2: Iteration: 10000 / 10000 [100%] (Sampling)
## Chain 2:
## Chain 2: Elapsed Time: 327.727 seconds (Warm-up)
## Chain 2:           337.567 seconds (Sampling)
## Chain 2:           665.294 seconds (Total)
## Chain 2:
##
## SAMPLING FOR MODEL 'continuous' NOW (CHAIN 3).
```

```
## Chain 3:
## Chain 3: Gradient evaluation took 0.0001 seconds
## Chain 3: 1000 transitions using 10 leapfrog steps per transition would take 1 seconds.
## Chain 3: Adjust your expectations accordingly!
## Chain 3:
## Chain 3:
## Chain 3: Iteration:    1 / 10000 [ 0%] (Warmup)
## Chain 3: Iteration: 1000 / 10000 [ 10%] (Warmup)
## Chain 3: Iteration: 2000 / 10000 [ 20%] (Warmup)
## Chain 3: Iteration: 3000 / 10000 [ 30%] (Warmup)
## Chain 3: Iteration: 4000 / 10000 [ 40%] (Warmup)
## Chain 3: Iteration: 5000 / 10000 [ 50%] (Warmup)
## Chain 3: Iteration: 5001 / 10000 [ 50%] (Sampling)
## Chain 3: Iteration: 6000 / 10000 [ 60%] (Sampling)
## Chain 3: Iteration: 7000 / 10000 [ 70%] (Sampling)
## Chain 3: Iteration: 8000 / 10000 [ 80%] (Sampling)
## Chain 3: Iteration: 9000 / 10000 [ 90%] (Sampling)
## Chain 3: Iteration: 10000 / 10000 [100%] (Sampling)
## Chain 3:
## Chain 3: Elapsed Time: 393.687 seconds (Warm-up)
## Chain 3:           239.748 seconds (Sampling)
## Chain 3:           633.435 seconds (Total)
## Chain 3:
##
## SAMPLING FOR MODEL 'continuous' NOW (CHAIN 4).
## Chain 4:
## Chain 4: Gradient evaluation took 0.000119 seconds
## Chain 4: 1000 transitions using 10 leapfrog steps per transition would take 1.19 seconds.
## Chain 4: Adjust your expectations accordingly!
## Chain 4:
## Chain 4:
## Chain 4: Iteration:    1 / 10000 [ 0%] (Warmup)
## Chain 4: Iteration: 1000 / 10000 [ 10%] (Warmup)
## Chain 4: Iteration: 2000 / 10000 [ 20%] (Warmup)
## Chain 4: Iteration: 3000 / 10000 [ 30%] (Warmup)
## Chain 4: Iteration: 4000 / 10000 [ 40%] (Warmup)
## Chain 4: Iteration: 5000 / 10000 [ 50%] (Warmup)
## Chain 4: Iteration: 5001 / 10000 [ 50%] (Sampling)
## Chain 4: Iteration: 6000 / 10000 [ 60%] (Sampling)
## Chain 4: Iteration: 7000 / 10000 [ 70%] (Sampling)
## Chain 4: Iteration: 8000 / 10000 [ 80%] (Sampling)
## Chain 4: Iteration: 9000 / 10000 [ 90%] (Sampling)
## Chain 4: Iteration: 10000 / 10000 [100%] (Sampling)
## Chain 4:
## Chain 4: Elapsed Time: 333.504 seconds (Warm-up)
## Chain 4:           297.088 seconds (Sampling)
## Chain 4:           630.592 seconds (Total)
## Chain 4:
```

DIAGNOSTICS

```
neff_ratio(wine_reg_model)
```

```
##                (Intercept)                points
##                0.30050                0.29950
##      b[(Intercept) country:Argentina]      b[points country:Argentina]
##                0.66650                0.67400
##      b[(Intercept) country:Australia]      b[points country:Australia]
##                0.61355                0.61650
##      b[(Intercept) country:Austria]        b[points country:Austria]
##                0.72895                0.73900
##      b[(Intercept) country:Chile]          b[points country:Chile]
##                0.68940                0.68965
##      b[(Intercept) country:France]         b[points country:France]
##                0.40045                0.39650
##      b[(Intercept) country:Germany]        b[points country:Germany]
##                0.66285                0.66370
##      b[(Intercept) country:Italy]          b[points country:Italy]
##                0.44205                0.43975
##      b[(Intercept) country:New_Zealand]     b[points country:New_Zealand]
##                1.03855                1.04565
##      b[(Intercept) country:Portugal]       b[points country:Portugal]
##                0.72290                0.72290
##      b[(Intercept) country:South_Africa]    b[points country:South_Africa]
##                1.06350                1.06915
##      b[(Intercept) country:Spain]          b[points country:Spain]
##                0.59200                0.59625
##      b[(Intercept) country:US]             b[points country:US]
##                0.34460                0.34310
##                sigma Sigma[country:(Intercept),(Intercept)]
##                1.02380                0.42940
##      Sigma[country:points,(Intercept)]      Sigma[country:points,points]
##                0.42780                0.42660
```

```
rhat(wine_reg_model)
```


```
##              (Intercept)                      points
##              1.0008012                      1.0008084
##      b[(Intercept) country:Argentina]      b[points country:Argentina]
##              1.0000986                      1.0000983
##      b[(Intercept) country:Australia]      b[points country:Australia]
##              1.0006402                      1.0006334
##      b[(Intercept) country:Austria]      b[points country:Austria]
##              1.0001205                      1.0001172
##      b[(Intercept) country:Chile]      b[points country:Chile]
##              1.0002377                      1.0002365
##      b[(Intercept) country:France]      b[points country:France]
##              1.0009156                      1.0009183
##      b[(Intercept) country:Germany]      b[points country:Germany]
##              0.9999520                      0.9999478
##      b[(Intercept) country:Italy]      b[points country:Italy]
##              1.0001172                      1.0001284
##      b[(Intercept) country:New_Zealand]      b[points country:New_Zealand]
##              1.0000847                      1.0000709
##      b[(Intercept) country:Portugal]      b[points country:Portugal]
##              1.0000758                      1.0000827
##      b[(Intercept) country:South_Africa]      b[points country:South_Africa]
##              1.0001118                      1.0001042
##      b[(Intercept) country:Spain]      b[points country:Spain]
##              1.0004439                      1.0004295
##      b[(Intercept) country:US]      b[points country:US]
##              1.0006993                      1.0007044
##              sigma Sigma[country:(Intercept),(Intercept)]
##              1.0002660                      1.0002212
##      Sigma[country:points,(Intercept)]      Sigma[country:points,points]
##              1.0002307                      1.0002394
```

```
tidy_sum<- tidy(wine_reg_model, effects= "fixed",
  conf.int = TRUE, conf.level = 0.90)
B0<- tidy_sum$estimate[1]
B1<- tidy_sum$estimate[2]
tidy_sum
```

```
## # A tibble: 2 × 5
##   term      estimate std.error conf.low conf.high
##   <chr>      <dbl>    <dbl>    <dbl>    <dbl>
## 1 (Intercept) -506.      75.4    -632.    -379.
## 2 points        6.10     0.864     4.64     7.54
```

```
wine_data %>%
  add_fitted_draws(wine_reg_model, n = 200, re_formula = NA)%>%
  ggplot(aes(x = points, y = price))+
  geom_line(aes(y = .value, group= .draw), alpha = 0.1) +
  geom_abline(intercept = B0, slope = B1, color = "blue")
```

```
## Warning in fitted_draws.default(model = model, newdata = newdata, ..., n = n): `fitted_draws`
## and `add_fitted_draws` are deprecated as their names were confusing.
## - Use [add_]epred_draws() to get the expectation of the posterior predictive.
## - Use [add_]linpred_draws() to get the distribution of the linear predictor.
## - For example, you used [add_]fitted_draws(..., scale = "response"), which
##   means you most likely want [add_]epred_draws(...).
## NOTE: When updating to the new functions, note that the `model` parameter is now
##   named `object` and the `n` parameter is now named `ndraws`.
```




```
# Get MCMC chains for the runner-specific intercepts & slopes
country_chains_2 <- wine_reg_model %>%
  spread_draws(`(Intercept)`, b[term, country], `points`) %>%
  pivot_wider(names_from = term, names_glue = "b_{term}",
              values_from = b) %>%
  mutate(country_intercept = `(Intercept)` + `b_(Intercept)`,
         country_points = points + b_points)
# Posterior medians of country-specific models
country_summaries_2 <- country_chains_2 %>%
  group_by(country) %>%
  summarize(country_intercept = median(country_intercept),
            country_points = median(country_points))
# Check it out
country_summaries_2
```



```
## # A tibble: 12 x 3
##   country          country_intercept country_points
##   <chr>              <dbl>          <dbl>
## 1 country:Argentina -411.           5.03
## 2 country:Australia -863.          10.2
## 3 country:Austria   -338.           4.12
## 4 country:Chile     -362.           4.42
## 5 country:France    -972.          11.4
## 6 country:Germany   -427.           5.16
## 7 country:Italy     -520.           6.27
## 8 country:New_Zealand -455.           5.49
## 9 country:Portugal  -367.           4.46
## 10 country:South_Africa -514.           6.19
## 11 country:Spain    -494.           5.98
## 12 country:US       -330.           4.14
```

```
pp_check(wine_reg_model)+
  labs(x = "rating")
```



```
set.seed(36501)
prediction_summary(model = wine_reg_model, data = wine_data)
```

```
##          mae mae_scaled within_50 within_95
## 1 11.72726   0.366392 0.7849794 0.9794239
```